

# EEE499 - Real-Time Embedded System Design

Modeling Real-Time Tasks



# Reference Model for Real-time Systems

- A reference model and consistent terminology let us reason about real-time systems
- A reference model is characterized by:
  - A **workload model** that describes the applications supported by the system
  - A **resource model** that describes the system resources available to the applications
  - **Algorithms** that define how the application system uses the resources at all times

# Timing Model

- We need to characterize jobs in order to schedule, manage and **reason about them with respect to time**

# Periodic Tasks Timing Model

A set of jobs that are executed at regular time intervals can be modelled as a periodic task.


$$T_i = (\phi_i, p_i, e_i, D_i)$$

Task      Phase      Period      Execution Time      Relative Deadline

# Periodic Tasks Timing Model

- The 4-tuple  $T_i = (\phi_i, p_i, e_i, D_i)$  refers to a periodic task  $T_i$  with
  - **phase  $\phi_i$** , is the release time  $r_{i,1}$  of the first job  $J_{i,1}$  in the task.
  - **period  $p_i$** , is the minimum length of all time intervals between release times of consecutive jobs
  - **execution time  $e_i$** , is the maximum execution time of all jobs in the periodic task
  - and **relative deadline  $D_i$**
  - Default phase of  $T_i$  is  $\phi_i = 0$ ,
  - Default relative deadline is the period  $D_i = p_i$

# Periodic Tasks Timing Model



# Precedence Constraints and Dependencies

- Jobs in a task may be constrained to execute in a particular order
  - Known as a precedence constraint
  - $j_i$  is a predecessor to another job  $j_k$  if  $j_k$  cannot begin execution until  $j_i$  completes its execution
    - Denoted by  $j_i < j_k$
    - $j_i$  is an immediate predecessor of  $j_k$  if there are no job such that  $j_i < j_j < j_k$
    - $j_i$  and  $j_k$  are independent when neither  $j_i < j_k$  nor  $j_k < j_i$
- A job with a precedence constraint **becomes ready** for execution when its release time has passed and all predecessors have completed

# Release Times

- The release time of a task is the time at which a task becomes available for execution
- Why *becomes* available?
- What if all tasks are always available for execution?



# Release Time


- We do not know exactly when the job will be released
  - Release time jitter due to
    - the granularity of clock tick (RTOS)
    - ISR time (depends on number of tasks and resources... why?)
- Even after it is ready to be executed, a task may still suffer **interference** from higher priority tasks
- And other interrupts...

# Release Time


- So we will have a release time variation
- $[r_i^-, r_i^+]$
- $r_i^-$  is the “earliest release time”
- $r_i^+$  is the “latest release time”
- For periodic tasks,  $r_i$  is the phase  $\phi_i$  *of the task*

# Release Time

- But what happens if we specify a release time that is not **effective**?



# Release Time




$$\underline{T_i(\phi_i, p_i, e_i, D_i)}$$

$$T_1(8, 10, 2, 16)$$

$$T_2(7, 15, 3, 15)$$

$$T_3(5, 12, 1, 12)$$




# Effective Release Time

- So the effective release time for task  $T_i$  is the maximum of its specified release time and the maximum effective release time of any of its predecessors + their execution times
  - Obviously this is recursive
  - We have not yet considered mutual exclusion constraints, just precedence due to data and control requirements

# Deadline

- The deadline of a task is the time at which a job must complete execution (within the clock tic)
- A deadline can be specified in two ways:
  - **Absolute Deadline:** Release Time plus a relative deadline
    - Limits our ability to reason about schedulability
  - **Relative Deadline:** Maximum Allowable Response-Time
    - Includes interference and blocking from other tasks... More to come.

# Deadline




$$\underline{T_i(\phi_i, p_i, e_i, D_i)}$$

$$T_1(8, 16, 2, 16)$$

$$T_2(7, 15, 3, 15)$$

$$T_3(5, 12, 1, 12)$$




# Effective Deadline

- In the same sense as for release times, the effective deadline is calculated from the precedence constraints.
  - More precisely from the successor constraints
- Effective deadline of a task is the minimum of the specified deadline for the task and the minimum effective deadline of all its successors minus their execution time



# Life of a task



# Processors and Resources

- Recall that a job executes on a special resource which we call a **processor**
- The job may also depend on some **resources**
- A processor,  $P$ , is an active component on which jobs are scheduled: i.e.
  - Threads scheduled on a CPU
  - Data scheduled on a transmission link
  - Read/write requests scheduled to a disk
  - Transactions scheduled on a database server

# Processors and Resources

- Processor (continued)
  - Each processor has a speed attribute that determines the rate of progress of a job
  - Two processors are of the same **type** if they are functionally identical and can be used interchangeably.
    - What would make two processors **heterogeneous**?

# Processors and Resources

- A resource,  $R$ , is a passive entity upon which jobs may depend: i.e.
  - Memory, sequence numbers, mutexes, database locks, mailboxes,...
  - Resources have different **types** and **sizes**, but do not have a speed attribute
  - Resources are usually **reusable**, and are not consumed by use

# Use of Resources

- If a system contains  $n$  types of resources it means:
  - There are  $n$  different types of **serially reusable** resources
  - There are one or more units of each type of resources, only one job can use each unit at once (Mutually exclusive access)
  - A job must obtain a unit of a needed resources, use it and release it

# Execution Time

- Perhaps one of the most difficult number to estimate is that of execution time.
- The execution time for job  $J_i$  varies in the interval  $[e_i^-, e_i^+]$  the interval depends on:
  - Conditional branches
  - Iterations in unbounded loops
  - Caches
- Without loss in generality we can have  $e_i = e_i^+$ 
  - The execution time for the job is therefore the **maximum execution time** ignoring the interval and lower bound


# Execution Time

- Needed to determine if deadlines can be met.
- What time do we choose?
  - The one from the last execution?
  - The one from the longest execution?
  - The average?
- How do we determine the execution time?

# Execution Time

How do we determine the execution time?

1. Analysis of the source code
2. Estimation from empirical evidence



[2]




# Release Time

- So we will have a release time variation
- $[r_i^-, r_i^+]$
- $r_i^-$  is the “earliest release time”
- $r_i^+$  is the “latest release time”
- For periodic tasks,  $r_i$  is the phase  $\varphi_i$  of the task

# Release and Response Time

- Release Time – The instant in time when a job becomes available for execution
  - May not be exact: Release time jitter in the interval  $[r_i^-, r_i^+]$
  - A job can be scheduled and executed at any time at or after its release time provided its resource dependency conditions are met
    - Including precedence constraints
- Response Time – the length of time from the release time of the job to the time instant when it completes
  - Not the same as execution time!

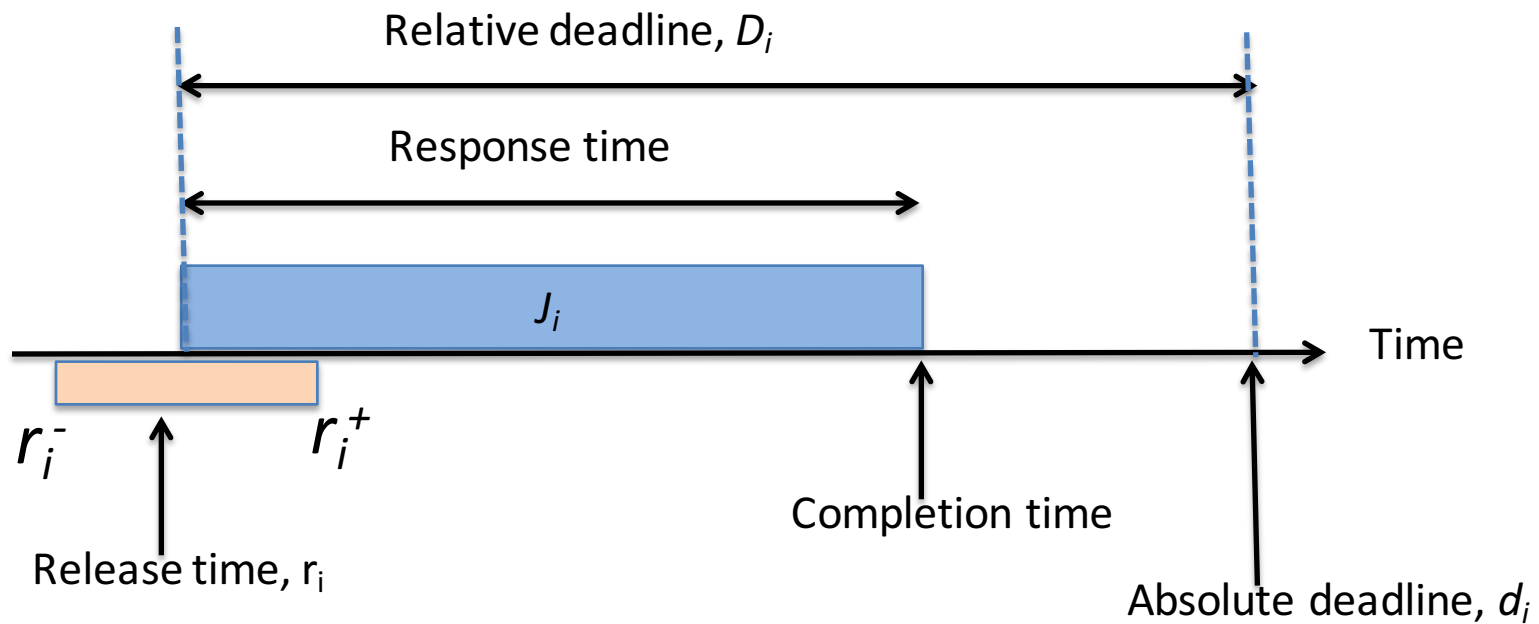
# Response Time



# Deadlines and Timing Constraints

- **Completion Time** – the instant at which a job completes execution
- **Relative deadline** – the maximum allowable job response time (**Schedulability analysis**)
- **Absolute deadline** – the instant of time by which a job is required to be completed
  - Absolute deadline = release time + relative deadline
  - Feasible interval for job  $J_i$  is the interval  $(r_i, d_i]$

# Response Time



# Aperiodic and Sporadic Jobs

- Many real-time systems are required to respond to unpredictable events.
- These are modelled as aperiodic or sporadic jobs
  - An aperiodic job has unpredictable release times
  - A sporadic job is a aperiodic job that has a hard deadline.
- Aperiodic jobs are always accepted.
- Sporadic jobs make the design of a hard real-time system impossible,
  - unless some bounds can be placed on their inter-arrival times and relative deadlines.
  - Based on the execution time and deadline of each newly arrived sporadic job, decide whether to accept or reject the job.

# References

[1] Liu, J. W. S. Real-Time Systems. Prentice Hall, 2000.

[2] Harder, D. W., Zarnett, J., Montaghani, V., Giannikouris, A. A practical introduction to real-time systems for undergraduate engineering, Version 0.2017.05.01. University of Waterloo, 2017. [Available online]